

Documentation Flow

Project - 3 {Regression Modelling Exercise}

- 1. Problem Statement->** We are provided with real time regression dataset and we have to perform EDA on that to find out and compare various machine learning algorithms and find which has the better rms value and accuracy and perform graphical visualization for better understanding of the data.
- 2. EDA->** After the data analysis part i.e. importing the dataset and removing null values. We used three machine learning algorithms and did comparative study namely- Logistic regression, Random Forest and XGBoost. We found out the accuracy, root mean squared error and mean squared error for each of them.

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
import scipy.stats as stats
```

```
In [2]: df = pd.read_csv("cancer_reg.csv", encoding = "ISO-8859-1")
```

```
In [3]: df.head()
```

```
Out[3]:
```

	avgAnnCount	avgDeathsPerYear	TARGET_deathRate	incidenceRate	medIncome	popEst2015	povertyPercent	studyPerCap
0	1397.0	469	164.9	489.8	61898	260131	11.2	499.748204
1	173.0	70	161.3	411.6	48127	43269	18.6	23.111234
2	102.0	50	174.7	349.7	49348	21026	14.6	47.560164

```
In [10]: df.drop('PctSomeCol18_24', axis=1, inplace=True)
```

```
In [11]: df.shape
```

```
Out[11]: (3047, 31)
```

```
In [12]: df['PctEmployed16_Over'].fillna(int(df['PctEmployed16_Over'].mean()), inplace=True)
```

```
In [13]: df.isnull().sum()
```

```
In [7]: df.describe()
```

```
Out[7]:
```

	avgAnnCount	avgDeathsPerYear	TARGET_deathRate	IncidenceRate	medIncome	popEst2015	povertyPercent	study
count	3047.000000	3047.000000	3047.000000	3047.000000	3047.000000	3.047000e+03	3047.000000	3047
mean	606.338544	185.965868	178.664063	448.268586	47063.281917	1.026374e+05	16.878175	155
std	1416.356223	504.134286	27.751511	54.560733	12040.090836	3.290592e+05	6.409087	529
min	6.000000	3.000000	59.700000	201.300000	22640.000000	8.270000e+02	3.200000	C
25%	76.000000	28.000000	161.200000	420.300000	38882.500000	1.168400e+04	12.150000	C
50%	171.000000	61.000000	178.100000	453.549422	45207.000000	2.664300e+04	15.900000	C
75%	518.000000	149.000000	195.200000	480.850000	52492.000000	6.867100e+04	20.400000	83
max	38150.000000	14010.000000	362.800000	1206.900000	125635.000000	1.017029e+07	47.400000	9762

8 rows × 32 columns

```
In [8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3047 entries, 0 to 3046
Data columns (total 32 columns):
#   Column                Non-Null Count  Dtype
---  -
0   avgAnnCount            3047 non-null   float64
1   avgDeathsPerYear       3047 non-null   int64
2   TARGET_deathRate       3047 non-null   float64
3   incidenceRate           3047 non-null   float64
4   medIncome              3047 non-null   int64
5   popEst2015             3047 non-null   int64
6   povertyPercent         3047 non-null   float64
7   studyPerCap            3047 non-null   float64
8   MedianAge              3047 non-null   float64
9   MedianAgeMale          3047 non-null   float64
10  MedianAgeFemale        3047 non-null   float64
11  AvgHouseholdSize       3047 non-null   float64
12  PercentMarried         3047 non-null   float64
13  PctNoHS18_24           3047 non-null   float64
14  PctHS18_24             3047 non-null   float64
15  PctSomeCol18_24        762 non-null    float64
16  PctBachDeg18_24        3047 non-null   float64
```

- 3. REGRESSION MODELING-** A regression model provides a function that describes the relationship between one or more independent variables and a response, dependent, or target variable.
- 4. COMPARATIVE STUDY-** We have used four machine learning algorithms for comparative study – Random Forest, XgBoost, Linear regression where we have observed that the root mean squared error value for Linear regression is the least and highest for XgBoost and on calculating the accuracy for each of them, we found that the linear regression was 91.68% accurate, whereas random forest was 88.51% accurate and XgBoost was 91.84% accurate which shows that XgBoost is the most accurate model.

Linear Regression

```
In [68]: from sklearn.linear_model import LinearRegression  
regressor = LinearRegression()
```

```
In [69]: regressor.fit(x_train, y_train)
```

```
Out[69]: LinearRegression()
```

```
In [70]: print(regressor.intercept_)
```

```
166.99646131143567
```

```
In [71]: print(regressor.coef_)
```

```
In [74]: from sklearn.metrics import r2_score  
r2_score(y_test, y_pred)
```

```
Out[74]: 0.5589216096657172
```

```
In [76]: from sklearn.metrics import mean_absolute_error, mean_squared_error
```

```
In [77]: mae = mean_absolute_error(y_test, y_pred)  
mse = mean_squared_error(y_test, y_pred)  
rmse = np.sqrt(mse)
```

```
In [78]: print(f'Mean absolute error: {mae:.2f}')  
print(f'Mean squared error: {mse:.2f}')  
print(f'Root mean squared error: {rmse:.2f}')
```

```
Mean absolute error: 14.16  
Mean squared error: 348.00  
Root mean squared error: 18.65
```

```
In [79]: #Measuring accuracy on Testing Data  
print('Accuracy', 100 - (np.mean(np.abs((y_test - y_pred) / y_test)) * 100))
```

```
Accuracy 91.68213848482719
```

Random Forest

```
In [80]: #Random Forest  
from sklearn.ensemble import RandomForestRegressor  
regressor = RandomForestRegressor(n_estimators = 10, random_state = 0)  
regressor.fit(x_train, y_train)
```

```
Out[80]: RandomForestRegressor(n_estimators=10, random_state=0)
```

```
In [81]: y_pred = regressor.predict(x_test)
```

```
In [82]: print(y_pred)
```

```
In [98]: from sklearn.metrics import r2_score
r2_score(y_test,y_pred)
```

```
Out[98]: 0.09562466141787007
```

```
In [84]: mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
```

```
In [85]: print(f'Mean absolute error: {mae:.2f}')
print(f'Mean squared error: {mse:.2f}')
print(f'Root mean squared error: {rmse:.2f}')
```

```
Mean absolute error: 14.43
Mean squared error: 380.50
Root mean squared error: 19.51
```

```
In [99]: #Measuring accuracy on Testing Data
print('Accuracy',100- (np.mean(np.abs((y_test - y_pred) / y_test)) * 100))

Accuracy 88.51488277507097
```

XGBOOST ALGO

```
In [86]: pip install xgboost
```

```
Requirement already satisfied: xgboost in c:\users\lenovo\anaconda3\envs\firstenv\lib\site-packages (1.6.1)
Note: you may need to restart the kernel to use updated packages.
Requirement already satisfied: scipy in c:\users\lenovo\anaconda3\envs\firstenv\lib\site-packages (from xgboost) (1.7.3)
Requirement already satisfied: numpy in c:\users\lenovo\anaconda3\envs\firstenv\lib\site-packages (from xgboost) (1.21.6)
```

```
In [101]: #XGBoost Algo
import xgboost as xg

xgb_r = xg.XGBRegressor(objective='reg:linear',
                        n_estimators = 10, seed = 123)

# Fitting the model
xgb_r.fit(x_train, y_train)
```

```
In [102]: from sklearn.metrics import r2_score
r2_score(y_test,y_pred)
```

```
Out[102]: 0.49573294824358105
```

```
In [103]: mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
```

```
In [104]: print(f'Mean absolute error: {mae:.2f}')
print(f'Mean squared error: {mse:.2f}')
print(f'Root mean squared error: {rmse:.2f}')
```

```
Mean absolute error: 14.47
Mean squared error: 397.85
Root mean squared error: 19.95
```


```
In [105]: #Measuring accuracy on Testing Data
print('Accuracy',100- (np.mean(np.abs((y_test - y_pred) / y_test)) * 100))

Accuracy 91.84954643555429
```

5. INFERENCE-> The graphical visualization of the dataset and their inference.

```
In [117]: plt.figure(figsize=(30,30))
sns.heatmap(df.corr(),cbar=True,annot=True,cmap='mako_r')
```

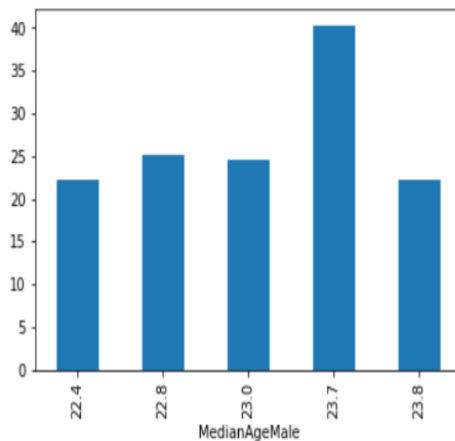
```
Out[117]: <AxesSubplot:>
```



Inference-> This graph shows the correlation values between the variables.


```
In [29]: df2.groupby(['MedianAgeMale'])['MedianAgeFemale'].sum().plot.bar()
```

```
Out[29]: <AxesSubplot:xlabel='MedianAgeMale'>
```



Inference-> This graph shows the relation between the MedianAgeMale and MedianAgeFemale


```
In [39]: sns.regplot(x='avgAnnCount', y='TARGET_deathRate', data=df);
```



Inference-> This graph shows the best fit of the data (best fit line).


```
In [34]: df4.groupby(['avgAnnCount'])['TARGET_deathRate'].sum().plot.bar()
```

```
Out[34]: <AxesSubplot: xlabel='avgAnnCount'>
```



Inference-> This graph shows the relation between the average count and target death rate.

```
In [43]: sns.jointplot(x='BirthRate', y='TARGET_deathRate', data=df, kind="reg");
```




Inference-> This graph shows the relation between the x and y i.e., how the dependent variable(y) varies with the independent variable(x).

```
In [48]: sns.distplot(df['TARGET_deathRate'])
```

```
C:\Users\Lenovo\anaconda3\envs\firstEnv\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: 'distplot' is a deprecated function and will be removed in a future version. Please adapt your code to use either 'displot' (a figure-level function with similar flexibility) or 'histplot' (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

```
Out[48]: <AxesSubplot:xlabel='TARGET_deathRate', ylabel='Density'>
```



Inference-> This graph shows the distribution plot of target death rate and tells us where the target values fall in a distribution.